

Elysian Truth Engine: Scaling Luxury with Certainty

Opal Agent – In-Instance Deliverables

- **Instruction Name (Main Agent):** Elysian Brand Director
 - **Specialized Agent Name(s):** Campaign Architect; Compliance Sentinel
 - **Workflow Name:** Elysian Content Workflow
 - **GitHub Repository:** [FDETH-TOOL](#)
-

a. Problem Identification:

- i. In the luxury sector, brand equity is built on precision and exclusivity. As marketing teams move towards fast, AI-driven content generation, they face a critical risk: **Hallucinations**. Standard LLMs often "guess" at sensitive data points like stock levels, rounded prices, or legal disclaimers.
- ii. For a brand like Elysian, an AI error regarding a £2,000 watch is a breach of customer trust and a potential legal liability. Marketing teams are currently forced to choose between **Speed of Creativity (AI)** and **Safety (Manual review)**. This solution eliminates that trade-off by tethering creativity to a programmatic *Source of Truth*.

b. Solution Overview:

- i. The **Elysian Truth Engine** is an agentic orchestration model that ensures content is audited against real-time enterprise data before reaching a customer. Utilizing a **Hub-and-Spoke architecture**, the system coordinates a "Creative" agent and a "Compliance" agent through a single governed workflow, verifying output against live inventory APIs and legal PDF standards.

c. Key Design Decisions & Implementation Notes:

- i. **Main agent's instructions and user interaction:** The **Elysian Brand Director** acts as the Global Orchestrator. Users interact via natural language (e.g., "Draft a campaign for our silver watch"). The agent identifies intent, enforces the brand's British luxury persona, and triggers the automated workflow.
- ii. **Role of Specialized Agent(s) and workflow integration:** Creation was decoupled from Compliance to prevent "instruction drift":
 1. **Campaign Architect:** Optimized for creative synthesis using custom tools and 'search_web'.
 2. **Compliance Sentinel:** A "Zero-Trust" auditor that cross-references drafts against the Truth Engine API and legal PDFs.

iii. Steps of the Agent Workflow and Opal tools used.

The workflow functions as a deterministic "Relay Race":

1. **Extraction:** Architect uses the **Custom Tool** to fetch authoritative stock/price data.
2. **Trend Synthesis:** Architect uses 'search_web' to enrich copy with market relevance.
3. **The Audit:** Sentinel performs a character-for-character check on legal footers and other guardrail information (micro-RAG).
4. **Final Asset:** Upon approval, Sentinel uses 'create_canvas' to generate the document for the user.

iv. Specific challenges and resolutions.

1. **Punctuation Hallucination:** LLMs often incorrectly "fix" legal footers. Remediation Logic was created in the Sentinel to auto-correct these (as opposed to just reject) while maintaining 100% PDF compliance.
2. **SKU Disambiguation:** Mapped "friendly" user names to technical IDs through a **Hardened Extraction Layer** to ensure resilient tool handshakes.

d. Self-Reflection:

i. How could this solution be extended or improved?

1. Future versions could include **Multi-Modal Validation** to ensure visual assets match verified SKU data, or a;
2. **Localization Agent** to adapt the "Elysian Truth" for global markets while maintaining regional legal compliance.

ii. What did you learn about Opal's capabilities?

Opal is an exceptionally powerful **Orchestration engine**. Effective design requires a **Separation of Concerns**, using specialized agents to manage high-entropy (creative) and low-entropy (auditing) tasks independently to create a "Self-Healing" system - this also allows for simple debugging of prompts and other potential hurdles.

It's also clear that Natural Language and Markup in Prompt Templates is absolutely essential for reasoning models like Gemini to really 'shine'.

In short: the tool is **awesome** and it's been a privilege to work with.